6.830 Lab 3: Query Optimization

Assigned: Wednesday, Mar 17, 2021

Due: Tuesday, Apr 6, 2021

In this lab, you will implement a query optimizer on top of SimpleDB.
The main tasks include implementing a selectivity estimation framework
and a cost-based optimizer. You have freedom as to exactly what you
implement, but we recommend using something similar to the Selinger
cost-based optimizer discussed in class (Lecture 9).

The remainder of this document describes what is involved in
adding optimizer support and provides a basic outline of how
you do so.

As with the previous lab, we recommend that you start as early as possible.

1. Getting started

You should begin with the code you submitted for Lab 2. (If you did not
submit code for Lab 2, or your solution didn't work properly, contact us to
discuss options.)

We have provided you with extra test cases as well
as source code files for this lab
that are not in the original code distribution you received. We again encourage you
to develop your own test suite in addition to the ones we have provided.

You will need to add these new files to your release. The easiest way
to do this is to change to your project directory (probably called simple-db-hw)
and pull from the master GitHub repository:

```
1    $ cd simple-db-hw
2    $ git pull upstream master
```

1.1. Implementation hints

We suggest exercises along this document to guide your implementation, but you may find that a different order makes
more sense for you. As before, we will grade your assignment by looking at your code and
verifying that you have passed the test for the ant targets `test` and
`systemtest` . See Section 3.4 for a complete discussion of grading and the tests you will need to pass.

Here's a rough outline of one way you might proceed with this lab. More details on these
steps are given in Section 2 below.

- Implement the methods in the `TableStats` class that allow
  it to estimate selectivities of filters and cost of
  scans, using histograms (skeleton provided for the `IntHistogram` class) or some
  other form of statistics of your devising.

- Implement the methods in the `JoinOptimizer` class that
  allow it to estimate the cost and selectivities of joins.
- Write the `orderJoins` method in `JoinOptimizer`. This method must produce
  an optimal ordering for a series of joins (likely using the
  Selinger algorithm), given statistics computed in the previous two steps.

2. Optimizer outline

Recall that the main idea of a cost-based optimizer is to:

- Use statistics about tables to estimate "costs" of different
  query plans. Typically, the cost of a plan is related to the cardinalities of
  (number of tuples produced by) intermediate joins and selections, as well as the
  selectivity of filter and join predicates.
- Use these statistics to order joins and selections in an
  optimal way, and to select the best implementation for join
  algorithms from amongst several alternatives.

In this lab, you will implement code to perform both of these
functions.

The optimizer will be invoked from `simpledb/Parser.java`. You may wish
to review the **lab 2 parser exercise**
before starting this lab. Briefly, if you have a catalog file
`catalog.txt` describing your tables, you can run the parser by
typing:

```
1    java -jar dist/simpledb.jar parser catalog.txt
```

When the Parser is invoked, it will compute statistics over all of the
tables (using statistics code you provide). When a query is issued,
the parser
will convert the query into a logical plan representation and then call
your query optimizer to generate an optimal plan.

2.1 Overall Optimizer Structure

Before getting started with the implementation, you need to understand the overall structure of the
SimpleDB
optimizer. The overall control flow of the SimpleDB modules of the parser and optimizer is shown in Figure
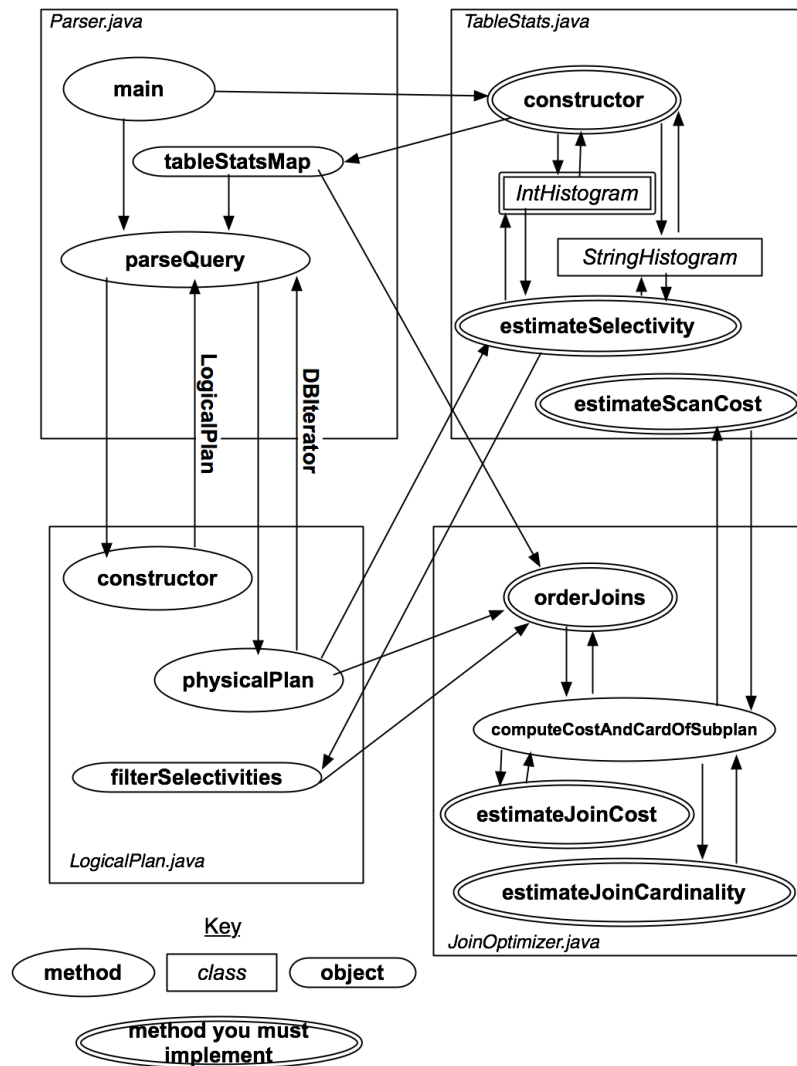1.

Figure 1: Diagram illustrating classes, methods, and objects used in the parser

The key at the bottom explains the symbols; you
will implement the components with double-borders. The classes and
methods will be explained in more detail in the text that follows (you may wish to refer back
to this diagram), but
the basic operation is as follows:

1. `Parser.java` constructs a set of table statistics (stored in the
   `statsMap` container) when it is initialized. It then waits for a
   query to be input, and calls the method `parseQuery` on that query.

2. `parseQuery` first constructs a `LogicalPlan` that
   represents the parsed query. `parseQuery` then calls the method `physicalPlan` on the
   `LogicalPlan` instance it has constructed. The `physicalPlan` method returns a
   `DBIterator` object that can be used to actually run the query.

In the exercises to come, you will implement the methods that help
`physicalPlan` devise an optimal plan.

2.2. Statistics Estimation

Accurately estimating plan cost is quite tricky. In this lab, we will focus only on the cost of sequences of joins and base table accesses. We won't worry about access method selection (since we only have one access method, table scans) or the costs of additional operators (like aggregates).

You are only required to consider left-deep plans for this lab. See Section 2.3 for a description of additional "bonus" optimizer features you might implement, including an approach for handling bushy plans.

2.2.1 Overall Plan Cost

We will write join plans of the form `p=t1 join t2 join ... tn`, which signifies a left deep join where t1 is the left-most join (deepest in the tree). Given a plan like `p`, its cost can be expressed as:

```
1    scancost(t1) + scancost(t2) + joincost(t1 join t2) +
2    scancost(t3) + joincost((t1 join t2) join t3) +
3    ...
```

Here, `scancost(t1)` is the I/O cost of scanning table t1, `joincost(t1,t2)` is the CPU cost to join t1 to t2. To make I/O and CPU cost comparable, typically a constant scaling factor is used, e.g.:

```
1    cost(predicate application) = 1
2    cost(pageScan) = SCALING_FACTOR x cost(predicate application)
```

For this lab, you can ignore the effects of caching (e.g., assume that every access to a table incurs the full cost of a scan) -- again, this is something you may add as an optional bonus extension to your lab in Section 2.3. Therefore, `scancost(t1)` is simply the number of pages in `t1 x SCALING_FACTOR`.

2.2.2 Join Cost

When using nested loops joins, recall that the cost of a join between two tables t1 and t2 (where t1 is the outer) is simply:

```
1    joincost(t1 join t2) = scancost(t1) + ntups(t1) x scancost(t2) //IO cost
2                         + ntups(t1) x ntups(t2)  //CPU cost
```

Here, `ntups(t1)` is the number of tuples in table t1.

2.2.3 Filter Selectivity

`ntups` can be directly computed for a base table by scanning that table. Estimating `ntups` for a table with one or more selection predicates over it can be trickier -- this is the *filter selectivity estimation* problem. Here's one

approach that you might use, based on computing a histogram over the values in the table:

- Compute the minimum and maximum values for every attribute in the table (by scanning it once).
- Construct a histogram for every attribute in the table. A simple approach is to use a fixed number of buckets $NumB$, with each bucket representing the number of records in a fixed range of the domain of the attribute of the histogram. For example, if a field f ranges from 1 to 100, and there are 10 buckets, then bucket 1 might contain the count of the number of records between 1 and 10, bucket 2 a count of the number of records between 11 and 20, and so on.
- Scan the table again, selecting out all of fields of all of the tuples and using them to populate the counts of the buckets in each histogram.
- To estimate the selectivity of an equality expression, $f=const$, compute the bucket that contains value $const$. Suppose the width (range of values) of the bucket is w, the height (number of tuples) is h, and the number of tuples in the table is $ntups$. Then, assuming values are uniformly distributed throughout the bucket, the selectivity of the expression is roughly $(h/w)/ntups$, since (h/w) represents the expected number of tuples in the bin with value $const$.
- To estimate the selectivity of a range expression $f>const$, compute the bucket b that $const$ is in, with width w_b and height h_b. Then, b contains a fraction $b\_f = h\_b/ntups$ of the total tuples. Assuming tuples are uniformly distributed throughout b, the fraction b_part of b that is $>const$ is $(b\_right - const)/w\_b$, where b_right is the right endpoint of b's bucket. Thus, bucket b contributes $(b\_f \times b\_part)$ selectivity to the predicate. In addition, buckets $b+1...NumB-1$ contribute all of their selectivity (which can be computed using a formula similar to b_f above). Summing the selectivity contributions of all the buckets will yield the overall selectivity of the expression. Figure 2 illustrates this process.
- Selectivity of expressions involving *less than* can be performed similar to the greater than case, looking at buckets down to 0.

Figure 2: Diagram illustrating the histograms you will implement in Lab 5

In the next two exercises, you will code to perform selectivity estimation of joins and filters.

Exercise 1: IntHistogram.java

You will need to implement some way to record table statistics for selectivity estimation. We have provided a skeleton class, `IntHistogram` that will do this. Our intent is that you calculate histograms using the bucket-based method described above, but you are free to use some other method so long as it provides reasonable selectivity estimates.

We have provided a class `StringHistogram` that uses `IntHistogram` to compute selecitivites for String predicates. You may modify `StringHistogram` if you want to implement a better estimator, though you should not need to in order to complete this lab.

After completing this exercise, you should be able to pass the `IntHistogramTest` unit test (you are not required to pass this test if you choose not to implement histogram-based selectivity estimation).

Exercise 2: TableStats.java

The class `TableStats` contains methods that compute the number of tuples and pages in a table and that estimate the selectivity of predicates over the fields of that table. The query parser we have created creates one instance of `TableStats` per

table, and passes these structures into your query optimizer (which you will need in later exercises).

You should fill in the following methods and classes in `TableStats` :

- Implement the `TableStats` constructor:
  Once you have
  implemented a method for tracking statistics such as histograms, you
  should implement the `TableStats` constructor, adding code
  to scan the table (possibly multiple times) to build the statistics
  you need.
- Implement `estimateSelectivity(int field, Predicate.Op op, Field constant)` : Using your statistics (e.g., an `IntHistogram`
  or `StringHistogram` depending on the type of the field), estimate
  the selectivity of predicate `field op constant` on the table.
- Implement `estimateScanCost()` : This method estimates the
  cost of sequentially scanning the file, given that the cost to read
  a page is `costPerPageIO` . You can assume that there are no
  seeks and that no pages are in the buffer pool. This method may
  use costs or sizes you computed in the constructor.
- Implement `estimateTableCardinality(double selectivityFactor)` : This method returns the number of tuples
  in the relation, given that a predicate with selectivity
  selectivityFactor is applied. This method may
  use costs or sizes you computed in the constructor.

You may wish to modify the constructor of `TableStats.java` to, for example, compute histograms over the fields as described above for purposes of selectivity estimation.

After completing these tasks you should be able to pass the unit tests in `TableStatsTest` .

2.2.4 Join Cardinality

Finally, observe that the cost for the join plan `p` above includes expressions of the form `joincost((t1 join t2) join t3)` . To evaluate this expression, you need some way to estimate the size (`ntups`) of `t1 join t2` . This *join cardinality estimation* problem is harder than the filter selectivity estimation problem. In this lab, you aren't required to do anything fancy for this, though one of the optional excercises in Section 2.4 includes a histogram-based method for join selectivity estimation.

While implementing your simple solution, you should keep in mind the following:

<!--

- **The following three paragraphs are different in this version of the lab.** *
  .-->
- For equality joins, when one of the attributes is a primary key, the number of tuples produced by the join cannot
  be larger than the cardinality of the non-primary key attribute.

- For equality joins when there is no primary key, it's hard to say much about what the size of the output is -- it could be the size of the product of the cardinalities of the tables (if both tables have the same value for all tuples) -- or it could be 0. It's fine to make up a simple heuristic (say, the size of the larger of the two tables).
- For range scans, it is similarly hard to say anything accurate about sizes. The size of the output should be proportional to the sizes of the inputs. It is fine to assume that a fixed fraction of the cross-product is emitted by range scans (say, 30%). In general, the cost of a range join should be larger than the cost of a non-primary key equality join of two tables of the same size.

Exercise 3: Join Cost Estimation

The class `JoinOptimizer.java` includes all of the methods for ordering and computing costs of joins. In this exercise, you will write the methods for estimating the selectivity and cost of a join, specifically:

- Implement
  `estimateJoinCost(LogicalJoinNode j, int card1, int card2, double cost1, double cost2)` : This method estimates the cost of join j, given that the left input is of cardinality card1, the right input of cardinality card2, that the cost to scan the left input is cost1, and that the cost to access the right input is card2. You can assume the join is an NL join, and apply the formula mentioned earlier.
- Implement `estimateJoinCardinality(LogicalJoinNode j, int card1, int card2, boolean t1pkey, boolean t2pkey)` : This method estimates the number of tuples output by join j, given that the left input is size card1, the right input is size card2, and the flags t1pkey and t2pkey that indicate whether the left and right (respectively) field is unique (a primary key).

After implementing these methods, you should be able to pass the unit tests `estimateJoinCostTest` and `estimateJoinCardinality` in `JoinOptimizerTest.java` .

2.3 Join Ordering

Now that you have implemented methods for estimating costs, you will implement the Selinger optimizer. For these methods, joins are expressed as a list of join nodes (e.g., predicates over two tables) as opposed to a list of relations to join as described in class.

Translating the algorithm given in lecture to the join node list form mentioned above, an outline in pseudocode would be:

```
1    1. j = set of join nodes
2    2. for (i in 1...|j|):
3    3.     for s in {all length i subsets of j}
4    4.        bestPlan = {}
5    5.        for s' in {all length d-1 subsets of s}
6    6.            subplan = optjoin(s')
7    7.            plan = best way to join (s-s') to subplan
8    8.            if (cost(plan) < cost(bestPlan))
9    9.                bestPlan = plan
10   10.       optjoin(s) = bestPlan
11   11. return optjoin(j)
```

To help you implement this algorithm, we have provided several classes and methods to assist you. First, the method `enumerateSubsets(List v, int size)` in `JoinOptimizer.java` will return a set of all of the subsets of `v` of size `size`. This method is VERY inefficient for large sets; you can earn extra credit by implementing a more efficient enumerator (hint: consider using an in-place generation algorithm and a lazy iterator (or stream) interface to avoid materializing the entire power set).

Second, we have provided the method:

```
1        private CostCard computeCostAndCardOfSubplan(Map<String, TableStats> stats,
2                                                     Map<String, Double>
    filterSelectivities,
3                                                     LogicalJoinNode joinToRemove,
4                                                     Set<LogicalJoinNode> joinSet,
5                                                     double bestCostSoFar,
6                                                     PlanCache pc)
```

Given a subset of joins (`joinSet`), and a join to remove from this set (`joinToRemove`), this method computes the best way to join `joinToRemove` to `joinSet - {joinToRemove}` . It returns this best method in a `CostCard` object, which includes the cost, cardinality, and best join ordering (as a list). `computeCostAndCardOfSubplan` may return null, if no plan can be found (because, for example, there is no left-deep join that is possible), or if the cost of all plans is greater than the `bestCostSoFar` argument. The method uses a cache of previous joins called `pc` (`optjoin` in the psuedocode above) to quickly lookup the fastest way to join `joinSet - {joinToRemove}` . The other arguments (`stats` and `filterSelectivities`) are passed into the `orderJoins` method that you must implement as a part of Exercise 4, and are explained below. This method essentially performs lines 6--8 of the psuedocode described earlier.

Third, we have provided the method:

```
1        private void printJoins(List<LogicalJoinNode> js,
2                               PlanCache pc,
3                               Map<String, TableStats> stats,
4                               Map<String, Double> selectivities)
```

This method can be used to display a graphical representation of a join plan (when the "explain" flag is set via
the "-explain" option to the optimizer, for example).

Fourth, we have provided a class `PlanCache` that can be used
to cache the best way to join a subset of the joins considered so far
in your implementation of Selinger (an instance of this class is
needed to use `computeCostAndCardOfSubplan`).

Exercise 4: Join Ordering

In `JoinOptimizer.java` , implement the method:

```
1     List<LogicalJoinNode> orderJoins(Map<String, TableStats> stats,
2                     Map<String, Double> filterSelectivities,
3                     boolean explain)
```

This method should operate on the `joins` class member,
returning a new List that specifies the order in which joins
should be done. Item 0 of this list indicates the left-most,
bottom-most join in a left-deep plan. Adjacent joins in the
returned list should share at least one field to ensure the plan
is left-deep. Here `stats` is an object that lets you find
the `TableStats` for a given table name that appears in the
`FROM` list of the query. `filterSelectivities`
allows you to find the selectivity of any predicates over a table;
it is guaranteed to have one entry per table name in the
`FROM` list. Finally, `explain` specifies that you
should output a representation of the join order for informational purposes.

You may wish to use the helper methods and classes described above to assist
in your implementation. Roughly, your implementation should follow
the psuedocode above, looping through subset sizes, subsets, and
sub-plans of subsets, calling `computeCostAndCardOfSubplan` and
building a `PlanCache` object that stores the minimal-cost
way to perform each subset join.

After implementing this method, you should be able to pass all the unit tests in
`JoinOptimizerTest` . You should also pass the system test
`QueryTest` .

2.4 Extra Credit

In this section, we describe several optional excercises that you may
implement for extra credit. These are less well defined than the
previous exercises but give you a chance to show off your mastery of
query optimization! Please clearly mark which ones you have chosen to complete in your report,
and briefly explain your implementation and present your results (benchmark numbers, experience reports,
etc.)

Bonus Exercises. Each of these bonuses is worth up to 5% extra credit:

- *Add code to perform more advanced join cardinality estimation* .
  Rather than using simple heuristics to estimate join cardinality,
  devise a more sophisticated algorithm.

  - One option is to use joint histograms between
    every pair of attributes a and b in every pair of tables $t1$ and $t2$.
    The idea is to create buckets of a , and for each bucket A of a , create a
    histogram of b values that co-occur with a values in A .
  - Another way to estimate the cardinality of a join is to assume that each value in the smaller table
    has a matching value in the larger table. Then the formula for the join selectivity would be:
    1/(*Max* (*num-distinct* (t1, column1), *num-distinct* (t2, column2))). Here, column1 and column2
    are the join attributes. The cardinality of the join is then the product of the cardinalities of $t1$ and
    $t2$ times the selectivity.
- *Improved subset iterator* . Our implementation of
  `enumerateSubsets` is quite inefficient, because it creates
  a large number of Java objects on each invocation.
  In this bonus exercise, you would improve the performance of
  `enumerateSubsets` so that your system could perform query
  optimization on plans with 20 or more joins (currently such plans
  takes minutes or hours to compute).

- *A cost model that accounts for caching* . The methods to
  estimate scan and join cost do not account for caching in the
  buffer pool. You should extend the cost model to account for
  caching effects. This is tricky because multiple joins are
  running simultaneously due to the iterator model, and so it may be
  hard to predict how much memory each will have access to using the
  simple buffer pool we have implemented in previous labs.

- *Improved join algorithms and algorithm selection* . Our
  current cost estimation and join operator selection algorithms
  (see `instantiateJoin()` in `JoinOptimizer.java`)
  only consider nested loops joins. Extend these methods to use one
  or more additional join algorithms (for example, some form of in
  memory hashing using a `HashMap`).

- *Bushy plans* . Improve the provided `orderJoins()` and other helper
  methods to generate bushy joins. Our query plan
  generation and visualization algorithms are perfectly capable of
  handling bushy plans; for example, if `orderJoins()`
  returns the list (t1 join t2 ; t3 join t4 ; t2 join t3), this
  will correspond to a bushy plan with the (t2 join t3) node at the top.

You have now completed this lab.
Good work!

3. Logistics

You must submit your code (see below) as well as a short (2 pages, maximum)
writeup describing your approach. This writeup should:

- Describe any design decisions you made, including methods for selectivity estimation,
  join ordering, as well as any of the bonus exercises you chose to implement and how

you implemented them (for each bonus exercise you may submit up to 1 additional page).
- Discuss and justify any changes you made to the API.
- Describe any missing or incomplete elements of your code.
- Describe how long you spent on the lab, and whether there was anything you found particularly difficult or confusing.
- Description of any extra credit implementation you have done.

3.1. Collaboration

This lab should be manageable for a single person, but if you prefer to work with a partner, this is also OK. Larger groups are not allowed. Please indicate clearly who you worked with, if anyone, on your writeup.

3.2. Submitting your assignment

We will be using gradescope to autograde all programming assignments. You should have all been invited to the class
instance; if not, please let us know and we can help you set up. You may submit your code multiple times before the
deadline; we will use the latest version as determined by gradescope. Place the write-up in a file called
lab3-writeup.txt with your submission. You also need to explicitly add any other files you create, such as new *.java files.

The easiest way to submit to gradescope is with `.zip` files containing your code. On Linux/MacOS, you can do so by
running the following command:

```
1    $ zip -r submission.zip src/ lab3-writeup.txt
```

3.3. Submitting a bug

SimpleDB is a relatively complex piece of code. It is very possible you are going to find bugs, inconsistencies, and bad, outdated, or incorrect documentation, etc.

We ask you, therefore, to do this lab with an adventurous mindset. Don't get mad if something is not clear, or even wrong; rather, try to figure it out
yourself or send us a friendly email.

Please submit (friendly!) bug reports to 6.830-staff@mit.edu .
When you do, please try to include:

- A description of the bug.
- A `.java` file we can drop in the `test/simpledb` directory, compile, and run.
- A `.txt` file with the data that reproduces the bug. We should be able to convert it to a `.dat` file using `HeapFileEncoder` .

You can also post on the class page on Piazza if you feel you have run into a bug.

3.4 Grading

75% of your grade will be based on whether or not your code passes the system test suite we will run over it. These tests will be a superset of the tests we have provided. Before handing in your code, you should make sure it produces no errors (passes all of the tests) from both `ant test` and `ant systemtest`.

Important: before testing, gradescope will replace your `build.xml`, `HeapFileEncoder.java` and the entire contents of the `test` directory with our version of these files. This means you cannot change the format
of `.dat` files! You should also be careful changing our APIs. You should test that your code compiles the unmodified tests.

You should get immediate feedback and error outputs for failed tests (if any) from gradescope after submission. The
score given will be your grade for the autograded portion of the assignment. An additional 25% of your grade will be
based on the quality of your writeup and our subjective evaluation of your code. This part will also be published on
gradescope after we finish grading your assignment.

We had a lot of fun designing this assignment, and we hope you enjoy hacking on it!